

Constraint Satisfaction Problem II

CSE 471 I: Artificial Intelligence


Md. Bakhtiar Hasan

Assistant Professor

Department of Computer Science and Engineering
Islamic University of Technology



k -Consistency




k -Consistency

■ Increasing degrees of consistency

- 1-Consistency (Node Consistency): Each single node's domain has a value which meets that node's unary constraints
- 2-Consistency (Arc Consistency): For each pair of nodes, any consistent assignment to one can be extended to the other
- k -Consistency: For each k nodes, any consistent assignment to $k - 1$ can be extended to the k^{th} node.


■ The higher the k , the more expensive to compute



Strong k -Consistency


- Also $k - 1, k - 2, \dots, 1$ consistent
- Claim: strong n -consistency means we can solve without backtracking!
 - Choose any assignment to any variable
 - Choose a new variable
 - By 2-consistency, there is a choice consistent with the first
 - Choose a new variable
 - By 3-consistency, there is a choice consistent with the first 2
 - ...
- Lots of middle ground between arc consistency and n -consistency!
(e.g. $k = 3$, called path consistency)

Problem Structure




Problem Structure

- Extreme case: independent subproblems
 - Example: Tasmania and mainland do not interact




Problem Structure

- Extreme case: independent subproblems
 - Example: Tasmania and mainland do not interact
- Independent subproblems are identifiable as connected components of constraint graph
 - Use DFS!




Problem Structure

- Extreme case: independent subproblems
 - Example: Tasmania and mainland do not interact
- Independent subproblems are identifiable as connected components of constraint graph
 - Use DFS!
- Suppose a graph of n variables can be broken into subproblems of only c variables:
 - Worst-case solution cost is $O((n/c)(d^c))$, linear in n
 - Compared to $O(d^n)$ for naïve backtracking
 - ▶ e.g., for $n = 80, d = 2, c = 20$
 - ▶ $2^{80} = 4$ billion years at 10 million nodes/sec
 - ▶ $(4)(2^{20}) = 0.4$ seconds at 10 million nodes/sec



Tree-Structured CSPs




- Theorem: if the constraint graph has no loops, the CSP can be solved in $O(nd^2)$ time
 - Compare to general CSPs, where worst-case time is $O(d^n)$
- Only one incoming arc per node
- Also applies to probabilistic reasoning

Tree-Structured CSPs

- Algorithm for tree-structured CSPs:


- Order: Choose a root variable, order variables so that parents precede children



Tree-Structured CSPs

- Algorithm for tree-structured CSPs:



- Order: Choose a root variable, order variables so that parents precede children



Tree-Structured CSPs

- Algorithm for tree-structured CSPs:



- Order: Choose a root variable, order variables so that parents precede children



Tree-Structured CSPs

■ Algorithm for tree-structured CSPs:



- Order: Choose a root variable, order variables so that parents precede children



Tree-Structured CSPs

■ Algorithm for tree-structured CSPs:


- Order: Choose a root variable, order variables so that parents precede children



Tree-Structured CSPs

- Algorithm for tree-structured CSPs:

- Order: Choose a root variable, order variables so that parents precede children




- Remove backward: For $i = n : 2$, apply $\text{RemoveInconsistent}(\text{Parent}(X_i), X_i)$

Tree-Structured CSPs

- Algorithm for tree-structured CSPs:

- Order: Choose a root variable, order variables so that parents precede children




- Remove backward: For $i = n : 2$, apply $\text{RemoveInconsistent}(\text{Parent}(X_i), X_i)$
- Assign forward: For $i = 1 : n$, assign X_i consistently with $\text{Parent}(X_i)$

Tree-Structured CSPs


■ Algorithm for tree-structured CSPs:

- Order: Choose a root variable, order variables so that parents precede children




- Remove backward: For $i = n : 2$, apply $\text{RemoveInconsistent}(\text{Parent}(X_i), X_i)$
 - Assign forward: For $i = 1 : n$, assign X_i consistently with $\text{Parent}(X_i)$
- ## ■ Runtime: $O(nd^2)$
- Go from tail to head, and then head to tail $\rightarrow O(n)$
 - Check pairs of values for consistency/assignment $\rightarrow O(d^2)$

Tree-Structured CSPs




- Claim #1: After backward pass, all root-to-leaf arcs are consistent

Tree-Structured CSPs




- Claim #1: After backward pass, all root-to-leaf arcs are consistent
 - Proof: Each $X \rightarrow Y$ was made consistent at one point and Y 's domain could not have been reduced thereafter (because Y 's children were processed before Y)

Tree-Structured CSPs




- Claim #1: After backward pass, all root-to-leaf arcs are consistent
 - Proof: Each $X \rightarrow Y$ was made consistent at one point and Y 's domain could not have been reduced thereafter (because Y 's children were processed before Y)
- Claim #2: If root-to-leaf arcs are consistent, forward assignment will not backtrack

Tree-Structured CSPs




- Claim #1: After backward pass, all root-to-leaf arcs are consistent
 - Proof: Each $X \rightarrow Y$ was made consistent at one point and Y 's domain could not have been reduced thereafter (because Y 's children were processed before Y)
- Claim #2: If root-to-leaf arcs are consistent, forward assignment will not backtrack
 - Proof: Induction on position

Tree-Structured CSPs



- Claim #1: After backward pass, all root-to-leaf arcs are consistent
 - Proof: Each $X \rightarrow Y$ was made consistent at one point and Y 's domain could not have been reduced thereafter (because Y 's children were processed before Y)
- Claim #2: If root-to-leaf arcs are consistent, forward assignment will not backtrack
 - Proof: Induction on position
- Why doesn't this algorithm work with cycles in the constraint graph?

Tree-Structured CSPs



- Claim #1: After backward pass, all root-to-leaf arcs are consistent
 - Proof: Each $X \rightarrow Y$ was made consistent at one point and Y 's domain could not have been reduced thereafter (because Y 's children were processed before Y)
- Claim #2: If root-to-leaf arcs are consistent, forward assignment will not backtrack
 - Proof: Induction on position
- Why doesn't this algorithm work with cycles in the constraint graph?
- Note: we'll see this basic idea again with Bayes' nets

Tree-Structured CSPs

function **TREE-CSP-SOLVER**(*csp*) returns a solution, or failure

local variables: X , set of variables

n , number of variables in X

$root$, any variable in X

D , set of values

$X \leftarrow \text{TOPOLOGICALSORT}(X, root)$

for $i \rightarrow n$ down to 2 do

MAKEARCCONSISTENT(**PARENT**(X_i), X_i)

 if no consistency then return failure


for $i \leftarrow 1$ to n

$X_i \leftarrow$ Any consistent value from D_i


 if no consistency then return failure

return X

Improving Structure




Nearly Tree-Structured CSPs




- Conditioning: instantiate a variable, prune its neighbors' domains
- Cutset conditioning: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree

Nearly Tree-Structured CSPs




- Conditioning: instantiate a variable, prune its neighbors' domains
- Cutset conditioning: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree
- Cutset size c gives runtime $O((d^c)(n - c)d^2)$, very fast for small c
 - Total number of instantiation: $O(d^c)$
 - Total number of remaining subproblems: $(n - c)$

Cutset Conditioning



Cutset Conditioning


Choose a cutset



Cutset Conditioning

Choose a cutset

Instantiate the cutset (all possible ways)




Cutset Conditioning

Choose a cutset

Instantiate the cutset (all possible ways)

Compute residual CSP for each assignment




Cutset Conditioning

Choose a cutset

Instantiate the cut-set (all possible ways)


Compute residual CSP for each assignment

Solve the residual CSPs (tree structured)




Cutset Quiz

- Find the smallest cutset for the graph below:



Iterative Improvement




Iterative Algorithms for CSPs

- Local search methods typically work with “complete” states, i.e., all variables assigned


Iterative Algorithms for CSPs

- Local search methods typically work with “complete” states, i.e., all variables assigned




Iterative Algorithms for CSPs

- Local search methods typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
 - Take an assignment with unsatisfied constraints
 - Operators *reassign* variable values
 - No fringe!




Iterative Algorithms for CSPs

- Local search methods typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
 - Take an assignment with unsatisfied constraints
 - Operators *reassign* variable values
 - No fringe!




Iterative Algorithms for CSPs

- Local search methods typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
 - Take an assignment with unsatisfied constraints
 - Operators *reassign* variable values
 - No fringe!
- Algorithm: While not solved
 - Variable selection: randomly select any conflicted variable
 - Value selection: min-conflicts heuristic:
 - ▶ Choose a value that violates the fewest constraints
 - ▶ i.e., hill climb with $h(n)$ = total number of violated constraints



Example: 4-Queens



- States: 4 queens in 4 columns ($4^4 = 256$ states)
- Operators: move queen in column
- Goal test: no attacks
- Evaluation: $c(n) = \text{number of attacks}$

Video: [5-queens-iterative-improvement](#)

Website: [complex - iterating improvement](#)


Performance of Min-Conflicts

- Given random initial state, can solve n-queens in almost constant time for arbitrary n with high probability (e.g., $n = 10,000,000$)!

Performance of Min-Conflicts

- Given random initial state, can solve n-queens in almost constant time for arbitrary n with high probability (e.g., $n = 10,000,000$)!
- The same appears to be true for any randomly-generated CSP except in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$




Local Search



Local Search

- Tree search keeps unexplored alternatives on the fringe (ensures completeness)
- Local search: improve a single option until you can't make it better (no fringe!)
- New successor function: local changes




- Generally much faster and more memory efficient (but incomplete and suboptimal)

Hill Climbing


- Simple, general idea
 - Start wherever
 - Repeat: move to the best neighboring state
 - If no neighbors better than current, quit
- What's bad about this approach?
 - Complete?
 - Optimal?
- What's good about it?



Hill Climbing Diagram



Hill Climbing Quiz



Starting from X, where do you end up?

Starting from Y, where do you end up?

Starting from Z, where do you end up?

Simulated Annealing

- Escape local maxima by allowing downhill moves → Make them rarer as time goes on

function **SIMULATED-ANNEALING**(*problem*, *schedule*) returns a solution state

inputs: *problem*, a problem

schedule, a mapping from time to “temperature”

local variables: *current*, a node

next, a node

T, a “temperature” controlling prob. of downward steps

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

for *t* \leftarrow 1 to ∞ do

T \leftarrow *schedule*[*t*]

if *T* = 0 then return *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ VALUE[*next*] - VALUE[*current*]

if $\Delta E > 0$ then *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$

Simulated Annealing

■ Theoretical guarantee

- Stationary distribution: $p(x) \propto e^{\frac{E(x)}{kT}}$
- If T decreased slowly enough, will converge to optimal state!



Simulated Annealing

- Theoretical guarantee
 - Stationary distribution: $p(x) \propto e^{\frac{E(x)}{kT}}$
 - If T decreased slowly enough, will converge to optimal state!
- Is this an interesting guarantee





Simulated Annealing

- Theoretical guarantee
 - Stationary distribution: $p(x) \propto e^{-\frac{E(x)}{kT}}$
 - If T decreased slowly enough, will converge to optimal state!
- Is this an interesting guarantee
- Sounds like magic, but reality is reality
 - The more downhill steps you need to escape a local optimum, the less likely you are to ever make them all in a row
 - People think hard about *ridge operators* which let you jump around the space in better ways




Genetic Algorithms



- Genetic algorithms use a natural selection metaphor
 - Keep best N hypotheses at each step (selection) based on a fitness function
 - Also have pairwise crossover operators, with optional mutation to give variety
- Possibly the most misunderstood, misapplied (and even maligned) technique around

Example: N-Queens



- Why does crossover make sense here?
- When wouldn't it make sense?
- What would mutation be?
- What would a good fitness function be?

Suggested Reading

- Russell & Norvig: Chapter 6.2-6.5