

Reinforcement Learning II


CSE 471 I: Artificial Intelligence

Md. Bakhtiar Hasan

Assistant Professor
Department of Computer Science and Engineering
Islamic University of Technology



Exploration vs. Exploitation



How to Explore?

- Several schemes for forcing exploration
 - Simplest: random actions (ϵ -greedy)
 - ▶ Every time step, flip a coin
 - ▶ With (small) probability ϵ , act randomly
 - ▶ With (large) probability $1 - \epsilon$, act on current policy



Videos: [q-bridge](#), [q-epsilon](#)

How to Explore?

- Several schemes for forcing exploration
 - Simplest: random actions (ϵ -greedy)
 - ▶ Every time step, flip a coin
 - ▶ With (small) probability ϵ , act randomly
 - ▶ With (large) probability $1 - \epsilon$, act on current policy
 - Problems with random actions?
 - ▶ You do eventually explore the space, but keep thrashing around once learning is done
 - ▶ One solution: lower ϵ over time
 - ▶ Another solution: exploration functions



Videos: [q-bridge](#), [q-epsilon](#)

Exploration Functions

- When to explore?
 - Random actions: explore a fixed amount
 - Better idea: explore areas whose badness is not (yet) established, eventually stop exploring

Video: [q-explorer-crawler](#)

Exploration Functions

- When to explore?
 - Random actions: explore a fixed amount
 - Better idea: explore areas whose badness is not (yet) established, eventually stop exploring



Video: [q-explorer-crawler](#)

Exploration Functions

- When to explore?
 - Random actions: explore a fixed amount
 - Better idea: explore areas whose badness is not (yet) established, eventually stop exploring
- Exploration function
 - Takes a value estimate u and a visit count n , and returns an optimistic utility, e.g.
$$f(u, n) = u + k/n$$



Video: [q-explorer-crawler](#)

Exploration Functions

- When to explore?
 - Random actions: explore a fixed amount
 - Better idea: explore areas whose badness is not (yet) established, eventually stop exploring
- Exploration function
 - Takes a value estimate u and a visit count n , and returns an optimistic utility, e.g.
$$f(u, n) = u + k/n$$

Regular Q-Update: $Q(s, a) \leftarrow_a R(s, a, s') + \gamma \max_{a'} Q(s', a')$



Video: [q-explorer-crawler](#)

Exploration Functions

- When to explore?
 - Random actions: explore a fixed amount
 - Better idea: explore areas whose badness is not (yet) established, eventually stop exploring

- Exploration function

- Takes a value estimate u and a visit count n , and returns an optimistic utility, e.g.

$$f(u, n) = u + k/n$$

Regular Q-Update: $Q(s, a) \leftarrow_a$

$$R(s, a, s') + \gamma \max_{a'} Q(s', a')$$

Modified Q-Update: $Q(s, a) \leftarrow_a$

$$R(s, a, s') + \gamma \max_{a'} f(Q(s', a'), N(s', a'))$$



Video: [q-explorer-crawler](#)

Exploration Functions

- When to explore?
 - Random actions: explore a fixed amount
 - Better idea: explore areas whose badness is not (yet) established, eventually stop exploring
- Exploration function
 - Takes a value estimate u and a visit count n , and returns an optimistic utility, e.g.
$$f(u, n) = u + k/n$$

Regular Q-Update: $Q(s, a) \leftarrow_a R(s, a, s') + \gamma \max_{a'} Q(s', a')$


Modified Q-Update: $Q(s, a) \leftarrow_a R(s, a, s') + \gamma \max_{a'} f(Q(s', a'), N(s', a'))$
 - Note: this propagates the “bonus” back to states that lead to unknown states as well!




Video: [q-explorer-crawler](#)

Regret

- Even if you learn the optimal policy, you still make mistakes along the way
- Regret is a measure of your total mistake cost: the difference between your (expected) rewards, including youthful suboptimality, and optimal (expected) rewards
- Minimizing regret goes beyond learning to be optimal – it requires optimally learning to be optimal
- Example: random exploration and exploration functions both end up optimal, but random exploration has higher regret



Approximate Q-Learning



Generalizing Across States

- Basic Q-Learning keeps a table of all q-values

Video: q-pacman, q-silent, q-tricky

Generalizing Across States

- Basic Q-Learning keeps a table of all q-values



Video: q-pacman, q-silent, q-tricky

Generalizing Across States

- Basic Q-Learning keeps a table of all q-values
- In realistic situations, we cannot possibly learn about every single state!
 - Too many states to visit them all in training
 - Too many states to hold the q-tables in memory



Video: [q-pacman](#), [q-silent](#), [q-tricky](#)

Generalizing Across States

- Basic Q-Learning keeps a table of all q-values
- In realistic situations, we cannot possibly learn about every single state!
 - Too many states to visit them all in training
 - Too many states to hold the q-tables in memory



Video: [q-pacman](#), [q-silent](#), [q-tricky](#)

Generalizing Across States

- Basic Q-Learning keeps a table of all q-values
- In realistic situations, we cannot possibly learn about every single state!
 - Too many states to visit them all in training
 - Too many states to hold the q-tables in memory
- Instead we want to generalize:
 - Learn about some small number of training states from experience
 - Generalize that experience to new, similar situations
 - This is a fundamental idea in machine learning, and we'll see it over and over again



Video: [q-pacman](#), [q-silent](#), [q-tricky](#)

Example: Pacman

Let's say we discover
through experience
that this state is bad:



Example: Pacman

Let's say we discover
through experience
that this state is bad:



In naïve q-learning,
we know nothing
about this state:



Example: Pacman


Let's say we discover
through experience
that this state is bad:



In naïve q-learning,
we know nothing
about this state:



Or even this one!



Feature-Based Representations

- Solution: describe a state using a vector of features (properties)
 - Features are functions from states to real numbers (often 0/1) that capture important properties of the state
 - Example features:
 - ▶ Distance to closest ghost
 - ▶ Distance to closest dot
 - ▶ Number of ghosts
 - ▶ $1/(\text{dist to dot})^2$
 - ▶ Is Pacman in a tunnel? (0/1)
 - Can also describe a q-state (s, a) with features (e.g. action moves closer to food)



Linear Value Functions


- Using a feature representation, we can write a Q-function for any state using a few weights:

$$\hat{Q}(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \cdots + w_n f_n(s, a)$$

- Advantage: our experience is summed up in a few powerful numbers
- Disadvantage: states may share features but actually be very different in value!

Approximate Q-Learning


$$\hat{Q}(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \cdots + w_n f_n(s, a)$$



Approximate Q-Learning

$$\hat{Q}(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \cdots + w_n f_n(s, a)$$

- Q-learning with linear Q-functions:




Approximate Q-Learning

$$\hat{Q}(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \cdots + w_n f_n(s, a)$$

- Q-learning with linear Q-functions:

Transition = (s, a, r, s')




Approximate Q-Learning

$$\hat{Q}(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \cdots + w_n f_n(s, a)$$

- Q-learning with linear Q-functions:

Transition = (s, a, r, s')

$$\text{Difference} = \left[r + \gamma \max_{a'} Q(s', a') \right] - \hat{Q}(s, a)$$



Approximate Q-Learning

$$\hat{Q}(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$


- Q-learning with linear Q-functions:

Transition = (s, a, r, s')

Difference = $\left[r + \gamma \max_{a'} Q(s', a') \right] - \hat{Q}(s, a)$

Exact Q's $\quad \underline{Q(s, a)} \leftarrow \underline{Q(s, a)} + \underline{\alpha [\text{difference}]}$

$$\underline{Q(s, a)} - \underline{\hat{Q}(s, a)}$$



Approximate Q-Learning

$$\hat{Q}(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \cdots + w_n f_n(s, a)$$

■ Q-learning with linear Q-functions:

Transition = (s, a, r, s')


Difference = $\left[r + \gamma \max_{a'} Q(s', a') \right] - \hat{Q}(s, a)$

Exact Q's $Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$

Approximate Q's $w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a)$

sample - $Q^*(s)$

$$W = W + \alpha [\text{difference}] F$$



Approximate Q-Learning

$$\hat{Q}(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \cdots + w_n f_n(s, a)$$

■ Q-learning with linear Q-functions:

Transition = (s, a, r, s')


Difference = $\left[r + \gamma \max_{a'} Q(s', a') \right] - \hat{Q}(s, a)$

Exact Q's $Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$

Approximate Q's $w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a)$

■ Intuitive interpretation:

- Adjust weights of active features
- E.g., if something unexpectedly bad happens, blame the features that were on: disprefer all states with that state's features



Approximate Q-Learning

$$\hat{Q}(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \cdots + w_n f_n(s, a)$$

■ Q-learning with linear Q-functions:

Transition = (s, a, r, s')

Difference = $\left[r + \gamma \max_{a'} Q(s', a') \right] - \hat{Q}(s, a)$


Exact Q's $Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$

Approximate Q's $w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a)$

■ Intuitive interpretation:

- Adjust weights of active features
- E.g., if something unexpectedly bad happens, blame the features that were on: disprefer all states with that state's features

■ Formal justification: online least squares




Example: Q-Pacman

$$\hat{Q}(s, a) = 4.0f_{DOT}(s, a) - 1.0f_{GST}(s, a)$$

Example: Q-Pacman


$$\hat{Q}(s, a) = 4.0f_{DOT}(s, a) - 1.0f_{GST}(s, a)$$



$$Q(s, \text{NORTH}) = +1$$

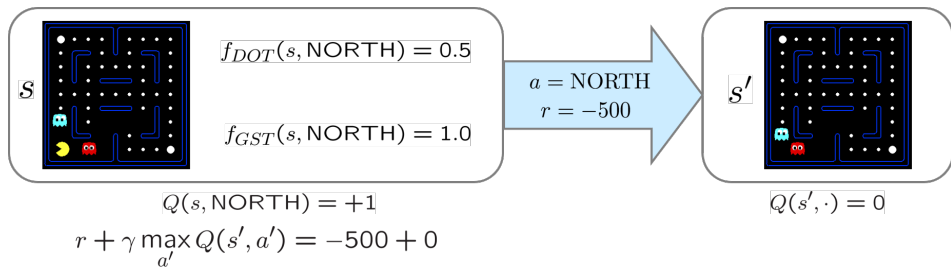
Example: Q-Pacman

$$\hat{Q}(s, a) = 4.0f_{DOT}(s, a) - 1.0f_{GST}(s, a)$$



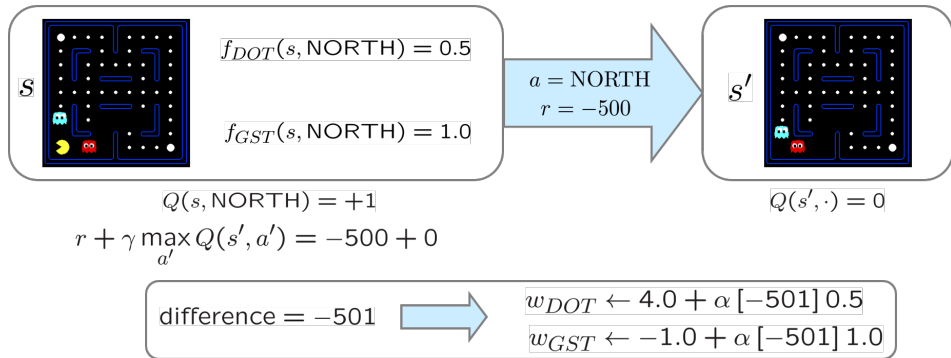
Example: Q-Pacman

$$\hat{Q}(s, a) = 4.0f_{DOT}(s, a) - 1.0f_{GST}(s, a)$$



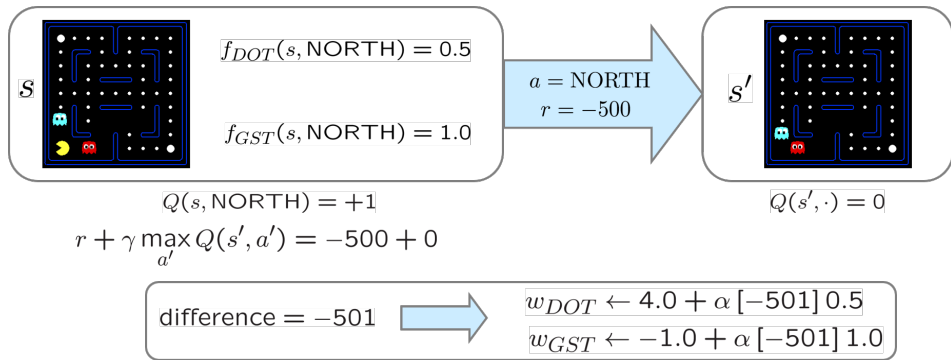
Example: Q-Pacman

$$\hat{Q}(s, a) = 4.0 f_{DOT}(s, a) - 1.0 f_{GST}(s, a)$$




Example: Q-Pacman

$$\hat{Q}(s, a) = 4.0f_{DOT}(s, a) - 1.0f_{GST}(s, a)$$




$$\hat{Q}(s, a) = 3.0f_{DOT}(s, a) - 3.0f_{GST}(s, a)$$

Q-Learning and Least Squares




Linear Approximation: Regression



Prediction:


$$\hat{y} = w_0 + w_1 f_1(x)$$

Linear Approximation: Regression



Prediction:


$$\hat{y} = w_0 + w_1 f_1(x)$$




Prediction:

$$\hat{y}_i = w_0 + w_1 f_1(x) + w_2 f_2(x)$$

Optimization: Least Squares




Optimization: Least Squares



$$\text{total error} = \sum_i (y_i - \hat{y}_i)^2$$

Optimization: Least Squares



$$\text{total error} = \sum_i (y_i - \hat{y}_i)^2 = \sum_i \left(y_i - \sum_k w_k f_k(x_i) \right)^2$$

Two blue arrows point from the y_i and \hat{y}_i terms in the equation to the corresponding terms in the scatter plot above.

Minimizing Error

Imagine we had only one point x , with features $f(x)$, target value y , and weights w :

Minimizing Error

Imagine we had only one point x , with features $f(x)$, target value y , and weights w :

$$error(w) = \frac{1}{2} \left(y - \sum_k (w_k f_k(x)) \right)^2$$

Minimizing Error

Imagine we had only one point x , with features $f(x)$, target value y , and weights w :

$$error(w) = \frac{1}{2} \left(y - \sum_k (w_k f_k(x)) \right)^2$$

$$\frac{\partial error(w)}{\partial w_m} = - \left(y - \sum_k w_k f_k(x) \right) f_m(x)$$

Minimizing Error

Imagine we had only one point x , with features $f(x)$, target value y , and weights w :

$$error(w) = \frac{1}{2} \left(y - \sum_k (w_k f_k(x)) \right)^2$$

$$\frac{\partial error(w)}{\partial w_m} = - \left(y - \sum_k w_k f_k(x) \right) f_m(x)$$

$$w_m \leftarrow w_m + \alpha \underbrace{\left(y - \sum_k (w_k f_k(x)) \right)}_{\text{exp}} \underbrace{f_m(x)}$$


Minimizing Error

Imagine we had only one point x , with features $f(x)$, target value y , and weights w :

$$error(w) = \frac{1}{2} \left(y - \sum_k (w_k f_k(x)) \right)^2$$

$$\frac{\partial error(w)}{\partial w_m} = - \left(y - \sum_k w_k f_k(x) \right) f_m(x)$$

$$w_m \leftarrow w_m + \alpha \left(y - \sum_k (w_k f_k(x)) \right) f_m(x)$$



Minimizing Error

Imagine we had only one point x , with features $f(x)$, target value y , and weights w :

$$error(w) = \frac{1}{2} \left(y - \sum_k (w_k f_k(x)) \right)^2$$


$$\frac{\partial error(w)}{\partial w_m} = - \left(y - \sum_k w_k f_k(x) \right) f_m(x)$$

$$w_m \leftarrow w_m + \alpha \left(y - \sum_k (w_k f_k(x)) \right) f_m(x)$$


Approximate q-update:

$$w_m \leftarrow w_m + \alpha \left[r + \gamma \max_a Q(s', a') - \underbrace{Q(s, a)}_{\text{same}} \right] f_m(s, a)$$


Q_-




Overfitting: Why Limiting Capacity Can Help




Overfitting: Why Limiting Capacity Can Help



Overfitting: Why Limiting Capacity Can Help



Policy Search



Policy Search

- Often the feature-based policies that work well aren't the ones that approximate V/Q best

Policy Search

- Often the feature-based policies that work well aren't the ones that approximate V/Q best
- Solution: Policy Search
 - Start with an initial linear value function or Q -function
 - Nudge each feature weight up and down and see if your policy is better than before

Policy Search


- Often the feature-based policies that work well aren't the ones that approximate V/Q best
- Solution: Policy Search
 - Start with an initial linear value function or Q -function
 - Nudge each feature weight up and down and see if your policy is better than before
- Problems
 - How do we tell the policy got better?
 - Need to run many sample episodes!
 - If there are a lot of features, this can be impractical

Policy Search

- Often the feature-based policies that work well aren't the ones that approximate V/Q best
- Solution: Policy Search
 - Start with an initial linear value function or Q -function
 - Nudge each feature weight up and down and see if your policy is better than before
- Problems
 - How do we tell the policy got better?
 - Need to run many sample episodes!
 - If there are a lot of features, this can be impractical
- Better methods exploit lookahead structure, sample wisely, change multiple parameters...

Conclusion

- We are done with Search and Planning!
- We have seen how AI methods can solve problems in:
 - Search
 - Constraint Satisfaction Problems
 - Games
 - Markov Decision Problems
 - Reinforcement Learning
- Next? Uncertainty and Learning



Suggested Reading

- Russell & Norvig: Chapter 21